

Contents

1. 1

(1) . . . . . 1

2. 1

2.a) . . . . . 1

2.d) . . . . . 2

3. 3

3.a) . . . . . 3

3.b) . . . . . 4

3.c) . . . . . 5

4. 12

4.a) . . . . . 12

4.b) . . . . . 14

4.c) . . . . . 15

1.

(1)

Wrong statements: C. E.  
## (2)  
Wrong statements: A. B. C.  
## (3)  
Wrong statements: C. ## (4)  
Wrong statements: A.  
## (5)  
Correct statements: B. D.  
## (6)  
Wrong statements: c. E.

2.

2.a)

$l(\theta) = \sum_{i=1}^n (\log(\theta-1) - \theta \log(X_i)) = n \log(\theta-1) - \theta \sum_{i=1}^n \log X_i \Rightarrow \frac{\partial}{\partial \theta} = \frac{n}{\theta-1} - \sum_{i=1}^n \log X_i = 0 \Rightarrow \hat{\theta}_n = \frac{n}{\sum \log X_i}$

## 2.b)

$$I(\theta) = - \int_1^\infty \left( \frac{\partial^2}{\partial \theta^2} \log((\theta-1)x^{-\theta}) \right) (\theta-1)x^{-\theta} dx = - \int_1^\infty \left( \frac{\partial}{\partial \theta} \frac{1}{(\theta-1)} - \log x \right) (\theta-1)x^{-\theta} dx = \int_1^\infty \left( \frac{1}{(\theta-1)} \right) x^{-\theta} dx = \int_1^\infty \frac{x^{-\theta}}{\theta-1} dx$$

$$(\hat{\theta}_n - \theta) \xrightarrow{d} N(0, \frac{1}{nI(\theta)}) \Rightarrow \lim_{n \rightarrow \infty} \frac{1}{nI(\theta)} = \lim_{n \rightarrow \infty} \frac{1}{n} (\theta-1)^2 = 0$$

## 2.c)

$$C_n = [\hat{\theta}_n - \frac{z_{\alpha/2}}{\sqrt{nI(\hat{\theta}_n)}}, \hat{\theta}_n + \frac{z_{\alpha/2}}{\sqrt{nI(\hat{\theta}_n)}}] = [\hat{\theta}_n - \frac{z_{0.025}}{\sqrt{n(\frac{1}{(\theta-1)^2})}}, \hat{\theta}_n + \frac{z_{0.025}}{\sqrt{n(\frac{1}{(\theta-1)^2})}}] = [\frac{1}{\log X_i} + 1 - \frac{1.96}{n(\frac{1}{(\theta-1)^2})}, \frac{1}{\log X_i} + 1 + \frac{1.96}{n(\frac{1}{(\theta-1)^2})}]$$

2.d)

$$F(x) = \int_1^\infty (\theta-1)x^{-\theta} dx = (-1+y) \left( \frac{x^{-y+1}}{-y+1} - \frac{1}{-y+1} \right) = (-1+2) \left( \frac{x^{-2+1}}{-2+1} - \frac{1}{-2+1} \right) = \left( \frac{x^{-1}}{-1} + 1 \right) = 1 - \frac{1}{x} \Rightarrow F^{-1}(x) = -\frac{1}{x-1} \quad (1)$$

```
set.seed(1003643587)
N <- 10000
n <- 100
t <- 2
invCDF <- function(u, t) {
  return(-1/(u-1))
}
```

confidence interval

```
lowerBound <- c()
upperBound <- c()
for (i in 1:N) {
  uni <- runif(100, 0, 1)
  samp <- c()
  for (j in 1:100) {
    samp <- append(samp, invCDF(uni[j], t), after = length(samp))
  }
  upperBound <- append(upperBound, 1 + 1/mean(log(samp)) + 1.96/(sqrt(n*mean(log(samp))^2)), after = length(upperBound))
  lowerBound <- append(lowerBound, 1 + 1/mean(log(samp)) - 1.96/(sqrt(n*mean(log(samp))^2)), after = length(lowerBound))
}
```

check for theta in confidence intervals

```
lowerBound = replace(lowerBound, lowerBound <= t, 1)
lowerBound = replace(lowerBound, lowerBound > t, 0)
upperBound = replace(upperBound, upperBound < t, 0)
upperBound = replace(upperBound, upperBound >= t, 1)

k <- 0
for (l in 1:N) {
  if(upperBound[l] == 1 & lowerBound[l] == 1) {
    k <- k + 1
  }
}
print(k/N)
```

## [1] 0.9483


As we can see, the frequency is 94.83% which is really close to 95%. Therefore we can claim that the constructed CI is effective

### 3.

#### 3.a)

```
set.seed(1003643587)
N <- 10
n <- c(10,100,1000,10000)


# function that samples from uniform
samples <- function(n) {
  samplesReplaced = runif(n, min = 0, max = 1)
  samplesReplaced = replace(samplesReplaced, samplesReplaced<0.5, -1)
  samplesReplaced = replace(samplesReplaced, samplesReplaced>= 0.5, 1)
  return(samplesReplaced)
}
```




As we can see from the plot above,  $\bar{X} - \mu$  is converging to 0 as  $n$  goes to infinity. for small  $n$ , we see that  $\bar{x} - \mu$  deviates from 0. However, as  $n$  increases, the plot gets increasingly closer to near 0.

3.b)


**epsilon = 0.5**



**epsilon = 0.1**



**epsilon = 0.05**



The plots above illustrates the law of large numbers. As  $n$  increases from 10 to 10000, we can see that the data points fall much closer to the mean( $\mu$ ), 0 for each of the three plots. We can also see that as epsilon becomes smaller, the data points for smaller  $n$ 's tend to be further away from  $\mu$ , however they still all will converge to  $\mu$  with large  $n$ .

**3.c)**

```
#install.packages("car")  
library(car)
```

```
## Warning: package 'car' was built under R version 4.0.3
```

```
## Loading required package: carData
```

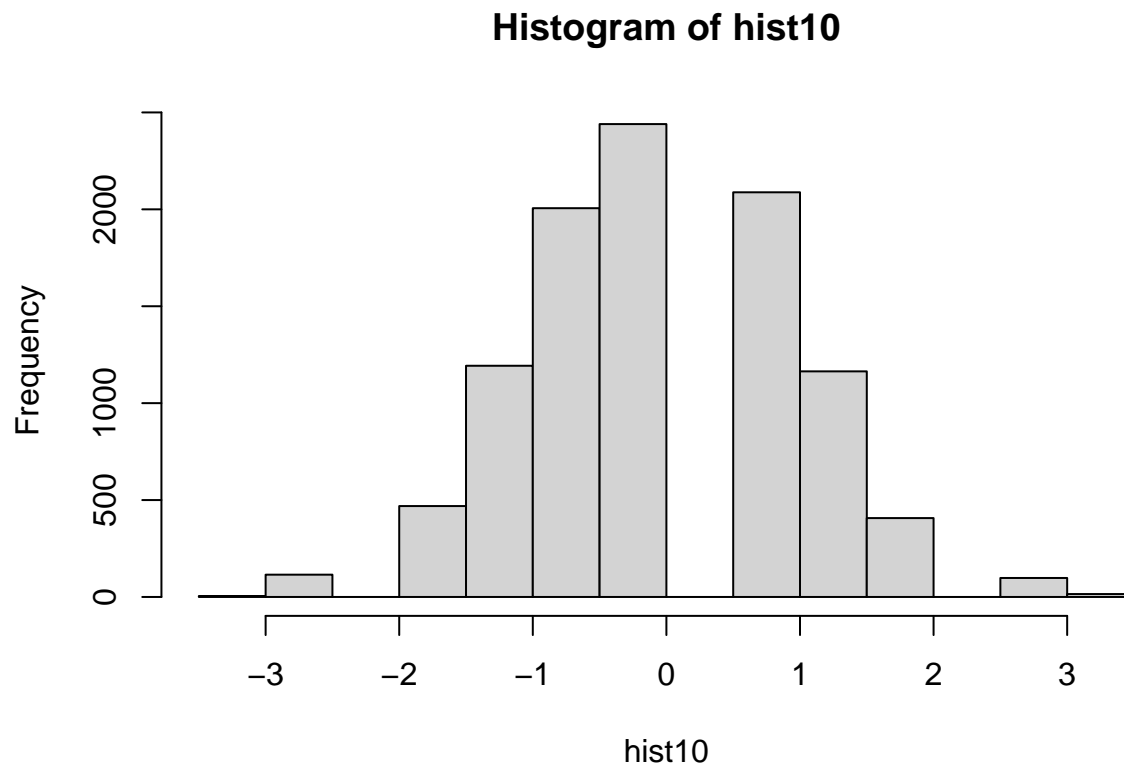
```
## Warning: package 'carData' was built under R version 4.0.3
```

```
set.seed(1003643587)  
mu <- 0  
# var(x) = E(X^2) = 1  
# sigma^2 = 1  
sigma <- 1  
mean10 <- c()  
mean100 <- c()  
mean1000 <- c()  
mean10000 <- c()  
for (i in 1:10000) {  
  mean10 <- append(mean10, sum(samples(10))/10)
```


```
mean100 <- append(mean100, sum(samples(100))/100)
mean1000 <- append(mean1000, sum(samples(1000))/1000)
mean10000 <- append(mean10000, sum(samples(10000))/10000)
}
```

n = 10

```
hist10 <- sqrt(10)*mean10/sigma
hist(hist10)
```




```
qqPlot(hist10)
```



```
## [1] 1286 2646
```


```
n = 1000
```

```
hist1000 <- sqrt(1000)*mean1000/sigma
hist(hist1000)
```



```
qqPlot(hist1000)
```






```
## [1] 2403 5974
```

```
n = 10000
```


```
hist10000 <- sqrt(10000)*mean10000/sigma
```

```
hist(hist10000)
```

**Histogram of hist10000**



```
qqPlot(hist10000)
```



```
## [1] 8914 2939
```


As we can see from the QQ plots and histograms, as  $n$  increases we can see that the data is more like standard normal. The histograms look more like normal as  $n$  increases and the qq plots are more linear with less outliers. These plots also show the Central Limit Theorem as we have  $\mu = E(x)$  and  $\sigma^2 = \text{var}(x)$ . As  $n$  approaches infinity,  $\sqrt{n}(\bar{x} - \mu)/\sigma$  approaches the standard normal in distribution.

```
## 3.d)
```

```
averag <- function(data, e) {
  data <- abs(data)
  data <- replace(data, data > e, 1)
  data <- replace(data, data <= e, 0)
  return(data)
}
```

```
logVals = c(log(10, base = 10), log(100, base = 10), log(1000, base = 10), log(10000, base = 10))
Y10 <- rnorm(10,0,1)
Y100 <- rnorm(100,0,1)
Y1000 <- rnorm(1000,0,1)
Y10000 <- rnorm(10000,0,1)
data10 <- sqrt(10)*mean10 - Y10
data10 <- averag(data10, 0.001)
data100 <- sqrt(100)*mean100 - Y100
data100 <- averag(data100, 0.001)
data1000 <- sqrt(1000)*mean1000 - Y1000
data1000 <- averag(data1000, 0.001)
data10000 <- sqrt(10000)*mean10000 - Y10000
```

```
data10000 <- averag(data10000, 0.001)
plotData <- c(mean(data10), mean(data100), mean(data1000), mean(data10000))
plot(logVals, plotData)
```



From the plot above, we see that it does not converge to  $Y$  in probability because even as  $n$  increases, the values do not necessarily approach 0. They still converge in distribution but this does not mean that all the values will be converging. As the standardized values of  $x$  approaches  $y$  in distribution, but not in probability.

4.

4.a)

```
#install.packages("bigmemory")
library(bigmemory)
```

```
## Warning: package 'bigmemory' was built under R version 4.0.3
```

```
X <- read.big.matrix("ratings.dat", type = "integer", col.names = c("UserID", "ProfileID", "Rating"))
head(X)
```

```
##      UserID ProfileID Rating
## [1,]  56669    39491      6
## [2,]  56919     8035     10
## [3,] 108853   102321     10
## [4,] 116784    52568      2
## [5,] 132748   220878     10
```

```
## [6,] 120139      29077      9
```


```
RatingAll <- X[, 3]  
C <- sum(RatingAll)/3000000
```

```
weighted.rank <- function(ProfileID) {  
  byProfile <- mwhich(X, "ProfileID", ProfileID, "eq")  
  newX <- X[byProfile, ]  
  byRating <- newX[, 3]  
  v <- length(byRating)  
  R <- sum(byRating)/v  
  m <- 4182  
  weightedRank <- ((v/(v+m))*R)+((m/(v+m))*C)  
  return(weightedRank)  
}  
weighted.rank(39491)
```

```
## [1] 5.977027
```

```
byUser <- mwhich(X, "UserID", 100, "eq")  
byUser <- X[byUser, ]  
byUser <- byUser[, 2]  
eachProfile <- lapply(byUser, weighted.rank)  
plotData <- unlist(eachProfile)  
hist(plotData, main = "weighted ranks of all profiles rated by UserID 100", xlab = "weighted rank")
```

### weighted ranks of all profiles rated by UserID 100



#### 4.b)

```
load("users.Rdata")
```

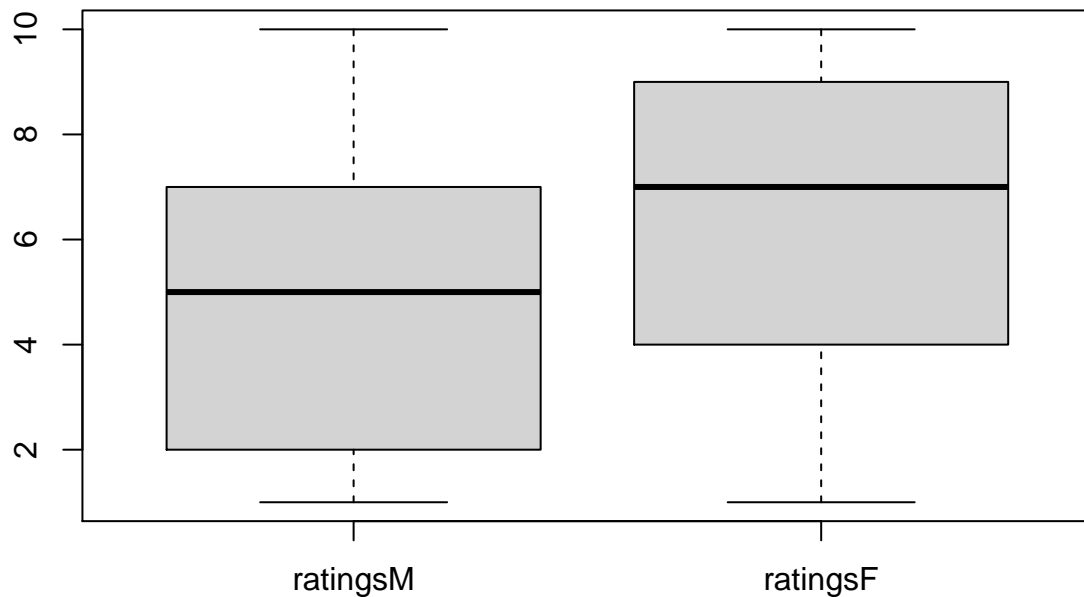
1) male users coming from New York State, State==New York

```
ratingsM <- c()
newYorkM <- User[which(User$Gender == 'M' & User$State == 'New York'), ]
newYorkM <- newYorkM[ , 1]
for(i in newYorkM) {
  byUser <- mwhich(X, "UserID", i, "eq")
  byUser <- X[byUser, ]
  if(identical(nrow(byUser), NULL)){
    byUser <- byUser[3]
    ratingsM <- append(ratingsM, byUser)
  } else {
    byUser <- byUser[ , 3]
    ratingsM <- append(ratingsM, byUser)
  }
}
```

2) female users coming from California, State==CA

```
ratingsF <- c()
CAF <- User[which(User$Gender == 'F' & User$State == 'CA'), ]
CAF <- CAF[ , 1]
for(i in CAF) {
  byUser <- mwhich(X, "UserID", i, "eq")
  byUser <- X[byUser, ]
  if(identical(nrow(byUser), NULL)){
    byUser <- byUser[3]
    ratingsF <- append(ratingsF, byUser)
  } else {
    byUser <- byUser[ , 3]
    ratingsF <- append(ratingsF, byUser)
  }
}
```

```
z <- c("ratingsM", "ratingsF")
histData <- lapply(z, get, envir=environment())
names(histData) <- z
boxplot(histData)
```



Above is a boxplot of ratings of Males from New York, against ratings of Females from CA

4.c)

```
#install.packages("biganalytics")
library(biganalytics)

## Warning: package 'biganalytics' was built under R version 4.0.3
## Loading required package: foreach
## Warning: package 'foreach' was built under R version 4.0.3
## Loading required package: biglm
## Warning: package 'biglm' was built under R version 4.0.3
## Loading required package: DBI
## Warning: package 'DBI' was built under R version 4.0.3

head(X)

##      UserID ProfileID Rating
## [1,]  56669    39491      6
## [2,]  56919     8035     10
## [3,] 108853   102321     10
## [4,] 116784    52568      2
## [5,] 132748   220878     10
## [6,] 120139    29077      9
```

```

N=3000000 # number of rating records Nu=135359 # maximum of UserID Np=220970 # maximum of
ProfileID user.rat=rep(0,Nu) # user.rat[i] denotes the sum of ratings given by user i user.num=rep(0,Nu) #
user.num[i] denotes the number of ratings given by user i profile.rat=rep(0,Np) # profile.rat[i] denotes the
sum of ratings given to profile i profile.num=rep(0,Np) # user.rat[i] denotes the number of ratings given to
profile i for (i in 1:N){ # In each iteration, we update the four arrays, i.e. user.rat, user.num, profile.rat,
profile.num, using one rating record. user.rat[X[i,'UserID']]=user.rat[X[i,'UserID']] + X[i,'Rating'] # The
matrix X here comes from the file 'ratings.dat' user.num[X[i,'UserID']]=user.num[X[i,'UserID']] + 1 pro-
file.rat[X[i,'ProfileID']]=profile.rat[X[i,'ProfileID']] + X[i,'Rating'] profile.num[X[i,'ProfileID']]=profile.num[X[i,'ProfileID']] + 1
if (i %% 10000 == 0) print(i/10000) } user.ave=user.rat/user.num profile.ave=profile.rat/profile.num
X1=big.matrix(nrow=nrow(X), ncol=ncol(X), type="double", dimnames=list(NULL, c('UsrAveRat','PrfAveRat','Rat'))))
X1[, 'Rat']=X[, 'Rating'] X1[, 'UsrAveRat']=user.ave[X[, 'UserID']] X1[, 'PrfAveRat']=profile.ave[X[, 'ProfileID']]
# X1 is the new data matrix we will work with in regression. user.ave profile.ave

```